

A Guided Introduction to Thinking *With* and *About* Object Oriented Programming

1 Introduction to Abstraction

The purpose of learning computer science, especially at the high school level, is not to memorize syntax. Rather, we are working to develop approaches to solving problems. Some of the best tools that we have to solve problems are programming languages.

Almost all programming languages help solve problems. Some languages are better suited to specific problems than other languages. There are many elements of a programming language that help solve problems. These elements can include built-in commands, data types, as well as the structure of the language itself.

When beginning a programming language, it is common to struggle over syntactical features such as commas, semi-colons, or indentation. Although, those elements must be correct for programs to work - they are *minor details* that simply support the true purpose of a programming language - **solving problems**. The structure of the programming language itself determines how we solve problems - namely, how we *de-construct* and think about problems.

When discussing programming languages, it is common to refer to *low level* and *high level* languages. By level, we are not referring to complexity but rather to the level of **abstraction**. Whereby **abstraction** refers to a general class - a broader collection of characteristics or qualities. For example, when discussing the color *red*, red does not refer to one specific color of one specific object, but rather refers to a color characteristic shared across objects that English speakers refer to as red.

With programming languages, *abstraction* can be seen in the amount of specificity commonly required and utilized across a program. For example, does each instance of *red* require a specific hex combination for each pixel on a screen to be manually specified or can the programmer more generally, simply refer to *red*?

A *low level* programming language such as **C** might require the programmer to be very specific and provide low level specifics such as actual memory bits of the computer. By contrast, a *high level* programming language such as **Processing**, **Python**, or **Java** may very well hide such details from the programmer - allowing the programmer to focus on more *abstract* i.e. 'big picture' ideas.

2 Object Oriented Programming

What is Object Oriented Programming (OOP)? Quite simply stated, OOP is programming structured around *objects* - conglomerations of qualities and characteristics - rather than always detailing specific interactions with the computer's memory. This *object oriented* approach is intended to simplify work for the programmer and improve programming efficiency by closer approximating real world complexity and interactions.

2.1 Generalization

A programmer can work to defining interactions across entire *classes* of objects rather than having to specify interactions for each specific instance.

But what does that mean????

Imagine - if every time someone bought a new pair of shoes they had to learn a specific system for tying the laces on that particular pair of shoes. That would be highly inefficient. Rather, we as a society have developed a few methods of tying shoes. Learning to tie one type of shoe readily transfers to tying other types of shoe laces. The laborious process of learning to tie a specific type of shoe lace does not have to be repeated each time. Similarly, imagine that a shoemaker would not want to invent and create a new system for tying shoes each time they made a pair of shoes. Rather, a system would be created for tying shoes and that system could be *inherited* by each pair of shoes that that shoemaker or other shoemakers could use.

2.2 Classes

Imagine further - someone has developed a good underlying design (or *template*) for shoes. In OOP, this can be referred to as a ***class***. In essence, a *class* is like a recipe that can be used as the basic for making many similar, but differentiated, *objects*. The shoemaker could create a *class* (template, recipe, etc.) for a generic shoe design. That class could then be used to mass produce however many shoes of that type (with whatever specified variations) that the shoemaker wants. In terms of OOP, the shoe template would be the *class* used to create individual shoe *objects*, whereby each pair of shoes would be considered an *instance* of **Class** Shoe.

As programmers, and problem solvers (e.g. think about Algebra), the goal is not to spend all our time focused on very specific problems. Rather, we attempt to represent and solve as many problems as possible at once. Sometimes, it may indeed be necessary to specify the behaviors for very particular interactions in very particular contexts. Cf., we may use particular greetings in particular situations (e.g. a specific handshake for a specific friend)- but before we develop specific greetings we develop a generalized greeting such as *Hi!* that can be broadly applicable.

But what does that mean???? In short, as programmers and people, it is not time efficient to develop specific greetings for each individual person

and take that approach to developing greetings for each possible random encounter. Someone could spend the better part of a month just planning out what they might say during a quick trip to the grocery store. Much more time effective is to plan a *generalized* i.e. *abstracted* approach and simply plan on saying “hello” to everyone and plan exceptions from there. For example, plan on saying “hello” to everyone except for people whose name is known - then say “hello `$personsname`.” This *abstracted* approach is the essence of OOP.

In Java, or other similarly *object oriented* programming languages, a class might simply be defined by something such as:

```
Class Person{ }
```

However, such a code by itself is useless - because this `Class Person` does not actually attribute any qualities to a person nor does it provide a way to create more Person objects i.e. people.

2.3 Data types

One of the biggest, most important concepts in programming, and perhaps in the modern world, is that there are different types of data. We use different types of data to answer different types of questions. For example, if someone asks “What type of cookies are those?” - the response “12” is not helpful or useful. Similarly, if someone asks “How many cookies do you have?” and is answered with “chocolate chip” - the response is neither helpful nor useful. A specific question requires a specific type of response.

In order to solve any problem, it is first important to understand what the problem is, the characteristics of the problem, and the likely characteristics of the solution. For example, if we are deciding on the types of cookies for a class party - some of the information we might need are numbers and type of cookies. *Number* and *Type* would be the data types we would use to discuss cookies. More cookie specific data types might include *sugar-type* or *leavening-agent*. As can be seen with the cookie example - certain data types make more sense based on the environment. In a cookie specific environment *sugar-type* and *leavening-agent* are very useful. However, in the design of cars they would basically be useless.

Data types in programming are those types of data that have been found to be most useful in programming i.e. computer related tasks. Therefore, each programming language has its own built in set of a data types. There is significant overlap - because there is significant overlap in the function of purpose of computer languages - also there is significant overlap because all computer languages must be designed to work on computers i.e. the actual specifics and architecture of processors determines the shape that a programming language takes - e.g. certain data types reflect the amount of space available in the computer’s working memory.

2.3.1 Data Types in Java

In Java, the built-in i.e. hard-coded data types are referred to as *primitives*. The *primitive* data types in Java include:

- Whole numbers
 - byte - a whole number (*integer*) between -128 and 127 (8 bits)
 - short - an *integer* between -32768 and 32767 (16 bits)
 - int - an *integer* between -2,147,483,648 and 2,147,483,647 (32 bits)
 - long - an *integer* between 9,223,372,036,854,775,808 and 9,223,372,036,854,755,807 (64 bits)
- Floating points (decimal numbers)
 - float - a decimal (*floating point*) (32 bits)
 - double - a decimal (*floating point*) (64 bits)
- Non numerical
 - char - a Unicode character e.g. one single letter, one single number, one single type of punctuation, etc.
 - boolean - true or false

See also the official Java documentation.

2.3.2 Strings

Another very common data type, that is **not** a *primitive* is the **String**. A **String** is simply an *array* (i.e. list) of characters. For example, `String name="John Doe"` contains all the characters and spaces for J-o-h-n-space-D-o-e. Because *Strings* are *objects*, they are able to perform actions (*methods*) not readily available to other primitives.

2.4 Using Data Types

In short, *classes* are recipes that create *objects* which do things (with *methods*) and contain *data*. There is perhaps no absolutely perfect way to organize data. However, some approaches to working with data are more useful than others. At the end of the day, programming, and indeed much of what we do in our lives, revolves around the organization and restructuring of data in new and useful ways.

Java is an *object oriented* programming language - which means that the language is structured to work with data as aspects of *objects*.

Earlier, the example of the empty `Class Person` was given, which was basically an empty recipe with no way of being made and no associated data. In order to do anything meaningful however, we will want to create *objects* that have data and can do things i.e. have *methods*.¹

2.5 First Worked Example - An *object*-first approach

In Java, basically everything occurs in relationship to some *Class*. Indeed, every Java program **must** have a *main class*.

2.5.1 The main class in Java

```
public class GenericProgramName {
    public static void main(String[] args) {
        // The magic starts here.
    }
}
```

The most important aspects to note (for now) in this example are:

- the curly braces `{}`
- the actual *main method* i.e.
`public static void main(String[] args) {}`
- the actual *class* name for the overarching program i.e.
`GenericProgramName`

In Java, units of code are enclosed in curly braces `{}`. The code between the first “{” and last “}” belongs together. Of course, this can, and often is, subdivided into further little bits. It may help to think of curly braces `{}` as boxes holding code. As with boxes, larger boxes can contain varied combinations of boxes within boxes - so too is true of curly braces cf.

{ { { { } } } }

Every Java program needs a *main* method to run. The majority of the program’s code may lie outside the *main* method - however, code execution (what is run) begins here.²

¹Note: many introductions to Java would include the typical “Hello World” example as the first program. However, to build a practiced, conceptual understanding of Java as an *object oriented* language - and to reinforce and better practice *object oriented* thinking - this text is written to emphasize an *objects* first approach.

²For more discussion of the *main* method in Java, see this Stack Overflow answer.

The `public class` declaration specifies the name of the overarching program *class* - consequently, the name of the file should match this *class declaration* i.e. `public class GenericProgramName` should be inside a file named `GenericProgramName.java`.

2.5.2 Most basic class with object in Java

The most basic program whereby we could create an object would look like this:

```
public class GenericProgramName {
    public static void main(String[] args) {
        Person billy = new Person();
    }
}

class Person {
    public Person() {
    }
}
```

There is the *public class* declaration with the name of the program (matching the file name). This contains the *main* method. In the *main* method all that is accomplished is an empty *Person object* named *billy* is created. Lastly, after the actual `public class GenericProgramName{...}` the `class Person` is created. However, `class Person` is basically empty - there are no data types associated with *Person* objects and no *methods* except for an empty **class constructor**: `public Person(){}`.

2.5.3 Using data types with a class in Java

When creating a *class Person*, it may be useful to provide some attributes. `class Person` is created here - there are no mandatory data types. It is at the discretion of the programmer to decide which data types may be most useful for `class Person`.³

Defining *Person* objects, the programmer may want to use common person related qualities such as: age, last name, first name, and whether the person is still a student. Consequently, the `class Person` will be created with the corresponding data types.

```
class Person {
    int age;
```

³The objects being created, as well as *object oriented* programming in general, are not structured solely to help the computer run best - but rather, the design is to help people utilize the computer most efficiently. The objects, their data, their data labels, their methods - are all approaches used by the programmer to make a program do what she decides. For the computer, everything boils down to 1's and 0's - everything else is to help the programmer use the computer (i.e. program the computer) most effectively.

```

String lastname;
String firstname;
boolean isStudent;
public Person() {
}
}

```

However, `class Person` still has the problem that age and name are not defined for Person objects. There are many possible solutions to this problem. Each solution has different advantages and disadvantages. As a programmer, we work not to find **THE** solution but rather to find *the better or most practical* solution to a given problem. As far as assigning common qualities to a Person object, we may want to start at the beginning - when an object is created i.e. when it is *instantiated*.

Objects, i.e. instances of a class, are created with the *class constructor*. Previously, the object *billy* was created with the line `Person billy = new Person();` where `new Person()` made a call to the *class constructor*.

2.5.4 Passing information to *constructors*

We will then create a class constructor to accept values for first name, last name, age, and whether someone is still a student. Values are passed *inside parentheses (...)*. The class constructor is the mechanism for creating new objects i.e. new *instances* of a class. Consequently, we must specify which data should be passed to a class constructor. The variable names used inside the class constructor are *only recognized inside the class constructor*. Consider:

```

public Person(String fn) {
    this.firstname = fn;
}
}

```

This class constructor allows us to create a person and specify their first name upon creation i.e. `Person billy = new Person("billy");` creates a new Person object with the value *billy* for the firstname variable. Note the use of `this.firstname` - the **this** is used to insure that the value specified is assigned to the *instance* being created. Here, the *this* is not completely necessary but is important for developing good coding habits. Further, the “.” specifies that the firstname is the value for “firstname” belonging to *this*. Elsewhere, the syntax `billy.lastname` could be used to directly refer to the object *billy*’s lastname value.⁴

In order to pass more values at the creation of each Person *instance*, we will need to specify the values to be passed. Continuing with in the same vein, we can create an additional constructor - one that looks almost identical but

⁴However, directly calling `billy.lastname` is not the preferred practice in Java. Rather, *getters* and *setters* should be used. [These are discussed later.]

accepts more information. In Java, having duplicate constructors (or methods) with the same name that only differ by the data they accept, is called *overloading*. Our new constructor that accepts first name, last name, and age is:

```
public Person(String fn, String ln, int a){
    this.firstname=fn;
    this.lastname=ln;
    this.age = a;
}
```

To use this constructor we would pass first name, last name, and age when creating a person object:

```
Person suzy = new Person("Suzy", "Jenkins", 12);
```

Notice: as with “billy” the names “Suzy” and “Jenkins” are in quotes. The quotation marks inform the *compiler* that we are sending a *String* and not a **variable**.

The program thus far might appear so:

```
public class GenericProgramName {
    public static void main(String[] args) {
        Person billy = new Person("billy");
        Person suzy = new Person("Suzy", "Jenkins", 12);
    }
}

class Person {
    int age;
    String lastname;
    String firstname;
    boolean isStudent;

    public Person() {
        //constructor one - no information is passed
    }

    public Person(String fn){
        this.firstname=fn;
        //constructor two - only first name is set
    }

    public Person(String fn, String ln, int a){
        this.firstname=fn;
        this.lastname=ln;
        this.age = a;
        //constructor three - first name, last name, and age are set
    }
}
```



```
    }
}
```

2.5.5 Doing stuff: Adding methods in Java

We have now created two *objects* - “billy” and “suzy” both *instances* of `class Person`. However, our objects cannot do anything as of yet. For our objects to do stuff, we will want to create *methods* that specify what our Person objects can do.

To create a method, we will need to set `access level/modifiers`, list the `return type`, give the method a `name`, specify any data to be passed (`parameters`), if any, and actually define what the method does (the `method body`).

First, we will create the method `sayHi()`. We will not specify an `access modifier`. The access level then defaults to *private-package* access. [The differences will be explained later.] Next, we specify the `return type` which will be `void` because we will not be passing any data back to the program. For the program name, we will choose `sayHi` - the actual action. We will not be passing any `parameters` - because we do not need additional information to say “hello.” For our actual command in the `method body` we will use the Java command to print a newline `System.out.println("Hello");`.

The resulting new method should be inside the Person class - because it is a method that belongs to `class Person`:

```
class Person {
    ...
    //The code has been trimmed to highlight the relevant bits.

    public Person(){
        ...
    }

    void sayHi(){
        System.out.println("Hello");
    }
}
```

The command itself - `System.out.println("Hello");` - (read backwards) consists of the `println()` command which prints a new line containing whatever text is between the parentheses i.e. (“Hello”). The `println()` command itself is a type of *output* i.e. information coming out of the

computer instead of going in. These *out* options for **output** belong to the very large number of built-in *System* commands available in Java.

The command itself is like a backwards address going from largest to smallest component - similar to if one were to write their address from largest to smallest unit as USA.Texas.San Antonio.Navarro.someaddress().

Now any Person object can use the method. For any method that belongs to a class, the object name is typed (here “suzy”) followed by a dot “.” and then the name of the method with any parameters inside the parentheses.⁵ The method is then called inside the main() method - after the object has been created⁶. The resulting code:

```
public static void main(String[] args) {
    ...
    suzy.sayHi();
}

int main() {
    printf("hello world");
    return 0;
}
```

⁵We are not passing any information to the method so our parentheses are empty.

⁶An object that has not yet been created cannot be called.